

Machine Learning (Lab support)

Regularization and feature selection

Abdelkrime Aries

*Laboratoire de la Communication dans les Systèmes Informatiques (LCSI)
École nationale Supérieure d'Informatique (ESI, ex. INI), Algiers, Algeria*

Academic year: 2024-2025





Attribution 4.0 International (CC BY 4.0)

<https://creativecommons.org/licenses/by/4.0/deed.en>

You are free to:

Share — copy and redistribute the material in any medium or format

Adapt — remix, transform, and build upon the material for any purpose, even commercially.

The licensor cannot revoke these freedoms as long as you follow the license terms.



Under the following terms:

Attribution — You must give appropriate credit, provide a link to the license, and indicate if changes were made. You may do so in any reasonable manner, but not in any way that suggests the licensor endorses you or your use.

No additional restrictions — You may not apply legal terms or technological measures that legally restrict others from doing anything the license permits.

Machine Learning (Lab support)

Regularization and feature selection: Introduction

- You saw in the lecture ...
 - the overfitting problem;
 - regularization by penalty (L2 as an example);
 - the different approaches to feature selection.
- In this lab support, we will present ...
 - some regularization approaches;
 - three penalty regularization techniques;
 - a reminder about feature selection approaches
 - detailing the techniques of each approach (with example from scikit-learn)

Machine Learning (Lab support)

Regularization and feature selection: Plan

1

Regularization

- L2 Loss
- L1 Loss
- ElasticNet

2

Feature selection

- Filter
- Embedded
- Wrapper

Regularization
Feature selection

L2 Loss
L1 Loss
ElasticNet

Section 1

Regularization

Regularization and feature selection

Regularization

- used to reduce overfitting
- can cause faster convergence
- **By augmentation:** add more data
 - automatic data generation
 - Ex. Generate new images by rotation
- **Early stopping:** use validation in the stopping decision
 - train on one dataset and use another for validation in each iteration
 - when the validation error increases, stop
- **By penalty:** add a penalty to the objective function
 - reduce model complexity
 - by adding another constraint on the parameters (penalty)

Regularization and feature selection

Regularization: By penalty

$$\left\{ \begin{array}{l} \min J_{\text{cost}}(\theta) \\ \wedge \\ \min J_{\text{complexity}}(\theta) \end{array} \right. \Rightarrow \min J_{\text{cost}}(\theta) + J_{\text{complexity}}(\theta)$$


- minimize the complexity by increasing the bias: θ_0 has no constraint, so the model gives it more importance during training
- the complexity penalty uses a hyper-parameter λ
 - if λ is too big, the model will be too simple (not dependent on attributes). So, **underfitting**.
 - if λ is too small, the model will be too complex (too dependent on attributes). So, **overfitting**.

Regularization and feature selection: Regularization

L2 Loss: Tikhonov regularization, L2 loss

$$J_{L2} = \frac{\lambda}{2M} \sum_{j=1}^N \theta_j^2$$

- θ_0 is not affected by regularization
- θ_j converges to 0, but does not equal 0
- $\lambda \rightarrow \infty \Rightarrow \theta_j \rightarrow 0$
- L2 tries to pull the values of θ inside a sphere
- the size of the sphere is inversely proportional to λ




Regularization and feature selection: Regularization

L1 Loss: L1 loss

$$J_{L1} = \frac{\lambda}{M} \sum_{j=1}^N |\theta_j|$$

- least absolute shrinkage and selection operator
- θ_0 is not affected by regularization
- θ_j is canceled after a few iterations (it will last depending on the importance of its attribute)
- $it \rightarrow \infty \Rightarrow \theta_j = 0$
- L1 tries to pull the values of θ inside a cube
- the size of the cube is inversely proportional to λ



Regularization and feature selection: Regularization

L1 Loss: Optimization

- L1 is not differential in 0
- Several techniques have been proposed to find the solution path $L1$
 - Subgradient methods; ex. least-angle regression (LARS)
 - Coordinate descent
 - Proximal gradient methods (by operator **soft thresholding**)

$$S_{\frac{\lambda}{M}}(\theta_j) = \begin{cases} \theta - \frac{\lambda}{M} & \theta_j > \frac{\lambda}{M} \\ 0 & \theta_j \in [-\frac{\lambda}{M}, \frac{\lambda}{M}] \\ \theta + \frac{\lambda}{M} & \theta_j < -\frac{\lambda}{M} \end{cases}$$

- In case of the proximal gradient, the parameters are updated as follows:
- $\theta = S_{\frac{\lambda}{M}}(\theta - \alpha \frac{\partial}{\partial \theta} J_{\text{cost}}(\theta))$

Regularization and feature selection: Regularization

ElasticNet

$$J_{EN} = \frac{\lambda}{M} \sum_{j=1}^N \left(r|\theta_j| + \frac{(r-1)}{2} \theta_j^2 \right)$$

- θ_0 is not affected by regularization
- $r \in [0, 1]$ controls the percentage of L1 regularization

Section 2

Feature selection

Regularization and feature selection

Feature selection

- can reduce overfitting: the model can overfit over noise features
- can improve system accuracy
- reduce training time
- **Filter**: independently of the estimator, the best features are selected based on univariate statistical tests
- **Embedded**: the best features are selected during training
- **Wrapper**: the best features are selected for a given estimator before training

Regularization and feature selection: Feature selection

Filter: Score

| Output | Input | |
|---------------------------------|------------------------------|------------------------------|
| | Numerical | Categorical |
| Numerical (Regression) | Pearson | Mutual information |
| Categorical (Classification) | ANOVA, Mutual information | Chi-2, Mutual information |

- **Pearson**: `sklearn.feature_selection.f_regression`
- **ANOVA**: `sklearn.feature_selection.f_classif`
- **Chi2**: `sklearn.feature_selection.chi2`
- **Mutual information**: `feature_selection.mutual_info_classif`
- **Mutual information**: `feature_selection.mutual_info_regression`

Regularization and feature selection: Feature selection

Filter: Selection

- Based on previous scores, the most important features can be selected in several ways
- **Number:** number of features to select
 - `sklearn.feature_selection.SelectKBest`
- **Percentile:** percentile of highest scores.
 - `sklearn.feature_selection.SelectPercentile`
- **P-value:** max threshold of accepted p-values
 - `sklearn.feature_selection.SelectFpr`

Regularization and feature selection: Feature selection

Filter: One-Way ANOVA (logic)

- Given a feature A (M samples) with numerical values
- the values are split on sets A_j where j is a class among N classes. We call them **Treatments**
- an attribute is representative (well correlated) of the output classes, if ...
 - values of the same class have less variance (intra-class variance)
 - class values have more variance with the rest (inter-class variance)
 - SO**, the ratio (inter-class variance)/(intra-class variance) must be large
 - We call this ratio: **F-value** of ANOVA

Regularization and feature selection: Feature selection

Filter: One-Way ANOVA (math)

- Correlation factor: $CF = \frac{(\sum_{ij} A_{ij})^2}{M}$
- Total Sum of Squares: $TotalSS = \sum_{ij} A_{ij}^2 - CF$
- Treatment Sum of Squares: $TreatmentSS = \sum_j \frac{(\sum_i A_{ij})^2}{|A_j|} - CF$
- Error Sum of Squares: $ErrorSS = TotalSS - TreatmentSS$
- Mean of Squares Between: $MSB = \frac{TreatmentSS}{(N-1)}$
- Mean of Squares Within: $MSW = \frac{ErrorSS}{(M-N)}$
- $Fvalue = \frac{MSB}{MSW}$

Regularization and feature selection: Feature selection

Embedded

- **Decision trees**

- In each node, the best split feature is chosen.
- Certain features will not be considered.
- Finally, the tree will only use features with high separation capacity.

- **L1 regularization**

- It forces parameters to have small values.
- In L1, certain parameters are set to 0 after convergence.
- In this case, during the estimation the model does not take into account features with parameters of zero.

Regularization and feature selection: Feature selection

Wrapper

- use an estimator and try to find feature combination giving more performance `sklearn.feature_selection.SequentialFeatureSelector`
- Ascending (Forward selection)**
 - the initial set of features is empty
 - start by selecting a single feature that maximizes cross-validation
 - add more features in the same way until reaching the desired number
 - `SequentialFeatureSelector(direction="forward")`
- Backward elimination**
 - the initial set of features equals to N
 - start by eliminating a single feature which minimizes cross-validation
 - eliminate more features in the same way until reaching the desired number
 - `SequentialFeatureSelector(direction="backward")`

The end of the ~~world~~ presentation

Stop scrolling

...

It is the end!